


# Ensemble learning

Daniel Lawson University of Bristol

Lecture 05.2 (v2.0.1)



<https://www.tcw.com/Insights/2023/2023-06-23-Agency-MBS-Feature>



What wins at Kaggle? From a 2019 Twitter post

# Signposting

- ▶ In 05.1 we introduced key baseline classification methods
- ▶ Other core methods based on trees and forests and in Block 06.
- ▶ This lecture is not about specific methods, but how they can be combined using **ensemble learning** (also called **meta-methods**).

# Questions

- ▶ If we have two good models, can we combine them into a better model?
- ▶ If we have many bad models, can they be combined into a good model?
- ▶ Can we apply cross-validation and other tricks to improve performance?

# Theory and Practice

- ▶ All of the practical work is done in the Workshop.
- ▶ **Activity 1** of the Workshop is about making a standardized interface for classifiers.
- ▶ This is an essential first step for using classifiers as building blocks for more complex operations.
- ▶ It tends to be where the Python Package SciKit Learn shines over R, though there are still inconsistencies

# Broad approaches to combining learners

- ▶ The goal is “Meta Learning”, i.e. combining multiple “learners”
  - ▶ learner = machine learning algorithms that consume data and make predictions
  - ▶ generalisation of a classifier
- ▶ There are two main approaches:
- ▶ **Parallel** approaches:
  - ▶ Run independently
  - ▶ Exploits independence structure between learners
- ▶ **Sequential** approaches:
  - ▶ Run dependently
  - ▶ Exploits dependence structure between learners
  - ▶ i.e. focus on what the previous set of learners are bad at

# Core topics

- ▶ Bagging
- ▶ Boosting
- ▶ Stacking



# Bagging

- ▶ The algorithm for bagging is straightforward - simply taking the average of bootstrapped learners
- ▶ In parallel,  $B$  times:
  - ▶ Form data sample  $\{X\}_b$ , e.g. by sampling with replacement, or leaving out a random subset of data
  - ▶ Learn a classifier  $f_b(x)$
- ▶ Output: A “bagged” classifier  $f(x) = \frac{1}{B} \sum_{b=1} f_b(x)$

# Bagging comments

- ▶ Bagging reduces overfitting to the data, and therefore works well on complex classifiers
- ▶ Same rules for resampling apply as in statistics: e.g. it works well when you respect the correlation structure
- ▶ In theory under certain assumptions, the **distribution of bagged learners** give a distribution on:
  - ▶ “what I could have seen if I obtained new data”
  - ▶ From the same distribution I got my data
- ▶ Usually little reason not to try it in practice

# Boosting

- ▶ The general idea of **Boosting** is:
  - ▶ Build a classifier, predict the data
  - ▶ Treat the residuals as “new data”
  - ▶ Repeat
- ▶ Boosting sounds like it should work for arbitrary classifiers, but because of the iterative nature it is applied to simple classifiers.
- ▶ There are many boosting algorithms, amongst which are:
  - ▶ Majority vote (Early and weak)<sup>1</sup>
  - ▶ Adaboost<sup>2</sup> (Adaptive boosting - first game-changer)
  - ▶ xgboost<sup>3</sup> (exploits sparsity and gradients)
  - ▶ LightGBM<sup>4</sup> (parallelizable and efficient)

---

<sup>1</sup>Kearns M and Valiant L (1989). Symp. Theor. Comp. ACM. 21: 433-444

<sup>2</sup>Freund and Schapire in 1996

<sup>3</sup>Chen T and Carlos G (2016) KDD 2016.

<sup>4</sup>**Microsoft dev team**

# Boosted feature splitter

- ▶ A very simple way to use boosting is to allow classifiers only from **single features**:
- ▶ Initialise weights of each **data sample** (uniformly)
- ▶ For  $T$  iterations:
  - ▶ Normalise weights
  - ▶ Train a classifier on every feature individually
  - ▶ Choose the best classifier, i.e. feature
  - ▶ Update the data weights by upweighting correct decisions and downweighting wrong decisions
- ▶ The boosted classifier uses a weighted sum of the selected classifiers

# Adaboost

- ▶ **Given:**  $N$  data  $(x_1, y_1), \dots, (x_N, y_N); x_i \in \mathcal{X}, y_i \in \{-1, 1\}$
- ▶ Set data weights  $D_{t=1}(n) = 1/N$ . For  $t = 1 \dots T$ :
  - ▶ **Train  $M$  “weak” classifiers**  $h_{mt}(x_t) : \mathcal{X} \rightarrow \{-1, 1\} \in \mathcal{H}$
  - ▶ With weighted **prediction error**  
 $\epsilon_{mt} = \sum_{i=1}^N D_t(i)(h_{mt}(x_i) - y_i)/2$
  - ▶ Choose the best classifier  $h_t = \operatorname{argmin}_m \epsilon_{mt}(h_{mt})$  with error  $\epsilon_t$
  - ▶ Evaluate  $\alpha_t = \log([1 - \epsilon_t]/\epsilon_t)$
  - ▶ Update the **weights**:

$$D_{t+1}(i) = \frac{D_t(i) \exp(\alpha_t \mathcal{I}(y_i \neq h_t(x_i)))}{Z_t}$$

- ▶ Where  $Z_t$  re-normalises weights  $D_{t+1}$  to sum to 1.
- ▶ **Output:** Boosted classifier:

$$H(x) = \operatorname{sign} \left( \sum_{m=1}^M \theta_m h_m(x) \right)$$

## Boosting comments

- ▶  $\alpha$  grows (towards infinity) as  $\epsilon$  shrinks (towards zero)
- ▶ The weighting process is chosen to ensure that the sign operation ensures correct classification
- ▶ Boosting is **computed** as a “decision tree” describing which classifier to use
- ▶ But outputs a mixture solution!
- ▶ All information about previous decisions is encoded into the weights
- ▶ When there is no residual error left for a data point, its weight is set to zero
- ▶  $h$  can be thought of as “features” and  $\mathcal{H} = \{h(x)\}$  can be large or infinite.
- ▶ Implementations in practice usually restrict weak classifiers  $h$  to a single, simple class (e.g. decision tree, perceptron)
- ▶ Weak classifiers are often generated by subsetting features, e.g. one at a time

# Stacking

- ▶ Stacking is a different way to combine multiple weak learners. It is more appropriate to combining “good classifiers” to make a meta-classifier.
- ▶ In theory a stacked classifier will always outperform its constituents if implemented appropriately<sup>5</sup>.
  - ▶ Cross-validation and asymptotics are required for this guarantee but in practice many approaches work.

---

<sup>5</sup>van der Laan, M, Polley E, Hubbard A, “Super Learner” (2007) Statistical Applications in Genetics and Molecular Biology, Volume 6.

# Super learner

- ▶ **Set up** the ensemble:
  - ▶ Specify  $L$  base classifiers.
  - ▶ Specify a metalearning algorithm.
- ▶ **Train** the ensemble:
  - ▶ Train the  $L$  base algorithms on the  $N$  training data. Use  $k$ -fold cross-validation for these learners.
  - ▶ For the  $N \times L$  matrix of predictions. Form the “level one” data with this matrix and the raw data.
  - ▶ Train the metalearning algorithm on the “level one” data.
- ▶ **Predict** on new data:
  - ▶ Generate predictions from the base classifiers.
  - ▶ Feed those predictions into the meta-learner to generate the ensemble prediction.



# More Stacking

- ▶ Related approaches:
  - ▶ Run any number of classification algorithms
  - ▶ Use their predictions as features
  - ▶ Use the data in addition to the predictions
- ▶ Pass this new feature set to any classification algorithm
- ▶ In practice, the best algorithm will be the one that generalises best in the test dataset. Common techniques:
  - ▶ Majority vote: use the prediction that most classifiers choose
  - ▶ Regularisation
  - ▶ Boosting-like prediction combination

# Wrapup

- ▶ Key to high prediction accuracy are:
  - ▶ **Complexity**: Non-linearity helps dramatically
  - ▶ **Bias control**: Don't overfit
  - ▶ **Meta-learning**: Boosting and stacking are essential for the final few percent.

# Signposting

- ▶ Next Block: Random Forests and decision trees and more practice using classification.
- ▶ Next Lecture: The workshop Lecture going over Bagging, Stacking and Boosting in practice.
- ▶ **References:**
- ▶ Ensemble learning in general:
  - ▶ Vadim Smolyakov, MIT: ML-perspective on Ensemble Methods
  - ▶ Stacked Ensembles by H2O, a Commercial AI Company focussing on Deployable AI
  - ▶ StackExchange: Stacking vs Bagging vs Boosting
  - ▶ Super Learners: van der Laan, M, Polley E, Hubbard A, "Super Learner" (2007) Statistical Applications in Genetics and Molecular Biology, Volume 6.

# Signposting (2)

- ▶ **More References:**
- ▶ Boosting:
  - ▶ AdaBoost paper: Experiments with a New Boosting Algorithm Freund and Schapire (1996).
  - ▶ Explaining AdaBoost, Rob Schapire, Empirical Inference (2013) pp 37-52.
  - ▶ xgboost Chen T and Carlos G (2016) KDD 2016.
  - ▶ xgboost explained, a blog post about Didrik Nilsen's paper Tree Boosting With XGBoost: Why Does XGBoost Win "Every" Machine Learning Competition?