

Latent Dirichlet Allocation

Daniel Lawson — University of Bristol

Lecture 07.1.3 (v1.0.2)

Signposting

- ▶ This lecture combines knowledge from 7.1.1 on topic models with 7.1.2 on Bayesian Methodology to describe one of the most successful tools in natural language processing.
- ▶ In 7.2 we cover some practicalities.
- ▶ In the workshop we implement these models in practice.

ILOs

- ▶ ILO2 Be able to use and apply basic machine learning tools
- ▶ ILO3 Be able to make and report appropriate inferences from the results of applying basic tools to data

Beyond the bag of words

- ▶ The Bag-of-words is a **vector representation** of a set of documents.
 - ▶ i.e. a feature space embedding.
- ▶ But how can we use this? How do we **compare** documents?
 - ▶ We could perform **dimensionality reduction** via PCA,
 - ▶ **Distance metrics** such as Cosine Similarity,
 - ▶ etc.
- ▶ Or we can model the similarity. The most successful approach for this is **Latent Dirichlet Allocation** (LDA).

Modelling a Bag Of Words using Latent Dirichlet Allocation

- ▶ Each document is modelled as a **mixture** of topics,
- ▶ Each topic is modelled as a **distribution** over words,
- ▶ Some Bayesian modelling magic allows the documents to be a theoretically **infinite mixture**,
- ▶ With content from ¹ which also contains cyber examples (without data).

¹Topic Modeling and Latent Dirichlet Allocation: An Overview (Weifeng Li, Sagar Samtani and Hsinchun Chen)

LDA Motivation - The setup

Topics

gene 0.04
dna 0.02
genetic 0.01
...

life 0.02
evolve 0.01
organism 0.01
...

brain 0.04
neuron 0.02
nerve 0.01
...

data 0.02
number 0.02
computer 0.01
...


Documents

Seeking Life's Bare (Genetic) Necessities

COLD SPRING HARBOR, NEW YORK—How many **genes** does an **organism** need to **survive**? Last week at the genome meeting here,* two genome researchers with radically different approaches presented complementary views of the basic genes needed for **life**. One research team, using **computer** analyses to compare known **genomes**, concluded that today's **organisms** can be sustained with just 250 genes, and that the earliest life forms required a mere 128 **genes**. The other researcher mapped genes in a simple parasite and estimated that for this organism, 800 genes are plenty to do the job—but that anything short of 100 wouldn't be enough.

Although the numbers don't match precisely, those **predictions**

"are not all that far apart," especially in comparison to the 75,000 **genes** in the human genome, notes Siv Andersson, a geneticist at the University of Uppsala, Sweden, who arrived at the 800 number. But coming up with a consensus answer may be more than just a **simple numbers** game—particularly if more and more **genomes** are being deciphered and sequenced. "It may be a way of organizing any newly **sequenced genome**," explains Arcady Mushegian, a **computational** molecular biologist at the National Center for Biotechnology Information (NCBI) in Bethesda, Maryland. Comparing in




* Genome Mapping and Sequencing, Cold Spring Harbor, New York, May 8 to 12.

Stripping down. Computer analysis yields an estimate of the minimum modern and ancient genomes.

SCIENCE • VOL. 272 • 24 MAY 1996

Topic proportions and assignments



LDA Motivation - Data in Practice

Topics



Documents


Topic proportions and assignments

Seeking Life's Bare (Genetic) Necessities

COLD SPRING HARBOR, NEW YORK—How many genes does an organism need to survive? Last week at the genome meeting here,* two genome researchers with radically different approaches presented complementary views of the basic genes needed for life. One research team, using computer analyses to compare known genomes, concluded that today's organisms can be sustained with just 250 genes, and that the earliest life forms required a mere 128 genes. The other researcher mapped genes in a simple parasite and estimated that for this organism, 800 genes are plenty to do the job—but that anything short of 100 wouldn't be enough.

Although the numbers don't match precisely, those predictions

"are not all that far apart," especially in comparison to the 75,000 genes in the human genome notes Siv Andersson, a Uppsala University in Sweden, who arrived at the 800 number. But coming up with a consensus answer may be more than just a genetic numbers game. Particularly as more and more genomes are completely ordered and sequenced. "It may be a way of organizing any newly sequenced genome," explains Arcady Mushegian, a computational molecular biologist at the National Center for Biotechnology Information (NCBI) in Bethesda, Maryland. Comparing an




* Genome Mapping and Sequencing, Cold Spring Harbor, New York, May 8 to 12.

Stripping down. Computer analysis yields an estimate of the minimum modern and ancient genomes.

SCIENCE • VOL. 272 • 24 MAY 1996

LDA Motivation - Example

The resulting output from an LDA model would be sets of topics containing keywords which would then be manually labeled. On the left are the inferred topic proportions for the example article from the pervious figure.



"Genetics"

human
genome
dna
genetic
genes
sequence
gene
molecular
sequencing
map
information
genetics
mapping
project
sequences

"Evolution"

evolution
evolutionary
species
organisms
life
origin
biology
groups
phylogenetic
living
diversity
group
new
two
common

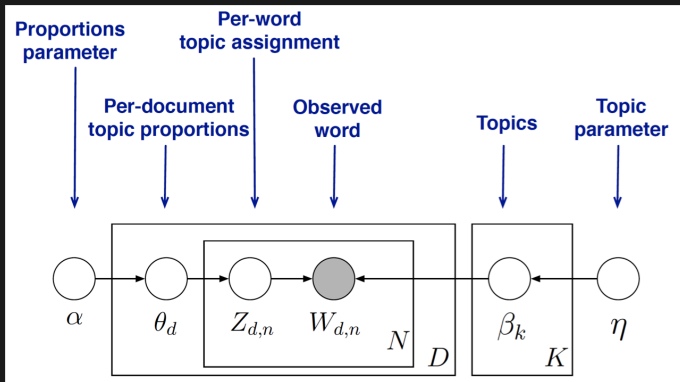
"Disease"

disease
host
bacteria
diseases
resistance
bacterial
new
strains
control
infectious
malaria
parasite
parasites
united
tuberculosis

"Computers"

computer
models
information
data
computers
system
network
systems
model
parallel
methods
networks
software
new
simulations

LDA Probabilistic Graphical Model



This is **plate notation** for **Bayesian Graphical Models**.

LDA Definition

- ▶ The overall **word distribution** is η , an N -vector.
- ▶ The overall **topic distribution** is α , a K -vector.
- ▶ Each **topic** k is described by a **word frequency** vector $\beta_k \sim \text{Dirichlet}(\eta)$.
- ▶ Each **document** d is described by a **topic frequency** vector $\theta_d \sim \text{Dirichlet}(\alpha)$.
- ▶ When generating word i from document d , we **generate a topic** $z_{di} \sim \text{Multinomial}(\theta_d)$.
- ▶ And then **generate a word** $w_{di} \sim \text{Multinomial}(\beta_d)$.

LDA properties

- ▶ Because it is a generative model, we can ask it to simulate documents.
- ▶ These approaches are embarrassing:
 - ▶ in the sense that if you simulate from the model, it generates garbage,
 - ▶ because words are independent.
- ▶ They should be thought of instead as keyword generators.
- ▶ This is extremely useful for a variety of text categorisation tasks.
- ▶ It can operate:
 - ▶ supervised (where we insist that some documents have pre-defined topic distributions) or
 - ▶ unsupervised (where nothing is assumed apriori about topics).

LDA implementation

- ▶ LDA implementations² use a conjugate model (Multinomial distribution is conjugate to the Dirichlet prior).
- ▶ It uses Variational Bayes to write the problem as an optimisation problem.

²Blei, David M., Andrew Y. Ng, and Michael I. Jordan. "Latent dirichlet allocation." Journal of machine Learning research 3, Jan (2003): 993-1022.

Further notes on LDA

- ▶ LDA models a matrix $Y = AX$, where:
 - ▶ Y is the data (N rows containing L word frequencies),
 - ▶ X are the topics (K rows containing L word frequencies) and
 - ▶ A is a mixture, (N rows containing K topics)
- ▶ This is a common problem called **matrix decomposition**.
- ▶ What makes LDA special is that **words are sparse**, meaning that there are many words but most words don't appear in most documents.
- ▶ You can run LDA on any problem of this type, but there are other approaches for dense data. (We return to **sparsity** later.)

Extensions

- ▶ We will not cover them, but if you work with document models you may want a more realistic model.
- ▶ Predictive text uses Markov Chains to predict $p(t(i)|d, t(i-1))$.
- ▶ Neural Networks generate arbitrary correlation structure, e.g.
 - ▶ Mathgen generates random papers,
 - ▶ Topic-RNN infers a topic model using a Neural Network.

Reflection

- ▶ What is a bag of words, conceptually?
- ▶ What are the advantages of LDA over Bag of words?
- ▶ And vice-versa?
- ▶ Could you use SVD on a bag of words?
- ▶ Why would we use either, when empirical accuracy of neural-network approaches is higher?

Signposting

- ▶ Next: Practicalities of text modelling.
- ▶ References:
 - ▶ Topic Modeling and Latent Dirichlet Allocation: An Overview (Weifeng Li, Sagar Samtani and Hsinchun Chen)
 - ▶ Blei, David M., Andrew Y. Ng, and Michael I. Jordan. “Latent dirichlet allocation”, Journal of machine Learning research 3.Jan (2003): 993-1022.