

Neural Network Architecture and Practicalities

Daniel Lawson — University of Bristol

Lecture 07.2 (v2.1.0)

Questions


- ▶ What are the most important types of neural net?
- ▶ What role does architecture have?

Some types of neural network

- ▶ Feed-forward
- ▶ Convolutional
- ▶ Recurrent
- ▶ Recursive
- ▶ Auto-encoders
- ▶ ...

Feed forward neural network

- This is the Neural Network that you know. It is acyclic.




Feed forward neural network

- ▶ The feed forward neural network is a **universal approximator**
- ▶ It can therefore be used as a **component** of a NN to compute **any** function $\mathbf{y} = f(\mathbf{x})$
- ▶ This can include:
 - ▶ **Likelihoods**, so making **probabilistic** predictions
 - ▶ **Derivatives**, (which are evaluated in the feed-forward step!)
 - ▶ And anything else we can imagine.
- ▶ Learning f can be complex, though many papers provide their network.
- ▶ Although all functions are approximable, not all behave nicely.
 - ▶ For example, densities seem hard to approximate whilst cumulative distribution functions behave better¹.

¹Chilinski and Silva **Neural Likelihoods via Cumulative Distribution Functions**


Convolutional neural network

- ▶ This is a feed-forward network that has carefully designed layers for constructing **known features**, such as local averaging.



- ▶ Choosing CNN architecture is **choosing a model**
- ▶ It should reflect known structure, e.g. locality, exchangeability, etc

Convolutional neural network




- ▶ CNNs are a core part of image processing²
- ▶ They scan an image, constructing **features**
- ▶ Different convolutions can create different features, including:
 - ▶ Larger objects
 - ▶ Edges
 - ▶ Presence/absence of either via max-pooling

²Albawi, Mohammed and Al-Zawi **Understanding of a convolutional neural network**


Recurrent Neural Network

- ▶ This is a network containing cycles, which allows for “memory” and potentially chaotic behavior.



- ▶ Training is hard; uses a special algorithm: “*causal recursive backpropagation*” which mitigates the disconnect between error and weights in standard algorithms. . .

Recurrent Neural Network for Point Processes




- ▶ An RNN acts as a “memory” for an arbitrary history³
- ▶ A CNN acts as a universal approximator to the CDF
- ▶ This is translated into the Likelihood of the data by back-propagation differentiation

³Omi, Ueda and Aihara Fully Neural Network based Model for General Temporal Point Processes

Recurrent Neural Network

- ▶ Recursive Neural Networks also exist, these allow cycles to previous layers. . .
- ▶ Alphago was an RNN. Alphago zero is better and used a “two-headed” architecture:
 - ▶ A **value network** that attributes values to board positions
 - ▶ A **policy network** that links board positions to actions that realise them
 - ▶ It is essentially making a giant decision tree, which is pruned to a manageable set by assigning values to states without seeing them through to outcomes.
- ▶ This is all beyond the scope of the course, but you might wish to examine how these work

Auto encoders

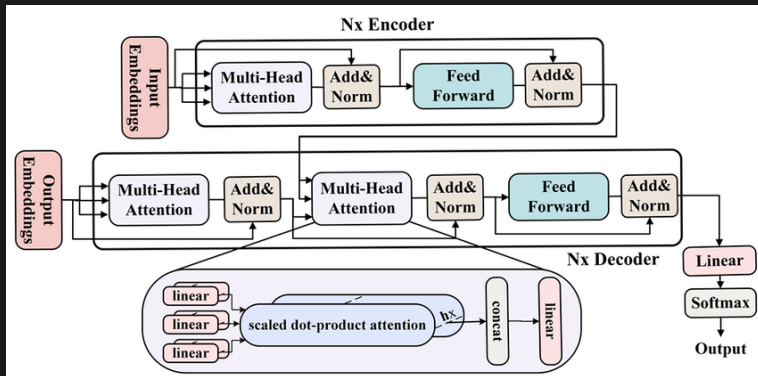


- ▶ Auto encoders provide a low-dimensional representation of the data
- ▶ They consist of separable parts, the encoder and the decoder
- ▶ They can be used for de-noising
- ▶ They are particularly useful when data are limited

LLMs, Foundation Models


- ▶ A Large Language Model (LLM) is the basis of modern Chatbots
- ▶ They are essentially very large transformers
- ▶ Trained on very large datasets
- ▶ To predict the next 'token'
- ▶ For a very long time!
- ▶ Attention is parallelizable learning, e.g. the fat cat sat on the mat
 - ▶ learns (the fat) \rightarrow cat and (the fat cat) \rightarrow sat simultaneously
 - ▶ solves the **vanishing gradient problem**, keeping context over long distances
- ▶ Foundation models are trained on more data types

Transformers



doi:10.48550/arXiv.2303.12914

Attention



From 'Attention is all you need' (2017) by A Vaswani · Cited by 137673

Finetuning

- ▶ In practice, **good neural networks are very large**
- ▶ Many problems contain very similar structure
- ▶ In practice you therefore want to **download someone else's model**
- ▶ You then **finetune** it to your task
- ▶ Best practice learns a 'low rank' approximation, e.g. LoRA (<https://github.com/microsoft/LoRA>)

Summary

- ▶ Neural Networks are possibly the most important development in AI.
- ▶ They provide universal approximation, allowing non-parametric approaches to wide problem sets
- ▶ Network design is critical, and still very much an art
- ▶ If you understand the building blocks just a little, you can access others' networks and potentially tweak them

Implementing Neural Networks

- ▶ Implementations are best thought of in two classes.
- ▶ **Simple networks** have a restricted architecture and can be deployed “out of the box” as a Machine Learning tool.
 - ▶ Examples include `sklearn.linear_model.Perceptron`, R’s `neuralnet` packages, etc
 - ▶ Often either shallow or very simple hidden layer structure
- ▶ **Deep networks** require a complex specification of architecture and significant computational optimisation, so are very large (and mercifully, open source) endeavours
 - ▶ This is the focus here.

Deep NN Implementations

- ▶ There are two main libraries for deep neural networks:
- ▶ **TensorFlow**, developed by Google Brain.
 - ▶ Well documented
 - ▶ Easier to use
 - ▶ Industry standard
 - ▶ **Tensorboard** visualisation is useful
- ▶ **PyTorch**, developed by Facebook.
 - ▶ Newer, less support
 - ▶ Dynamical coding paradigm: graph can remodel in the light of the data
 - ▶ Debugging is easier? As the code is compiled at runtime, like native python

Using implementations

- ▶ **Tensorflow** is a low-level language. You can interact with it through abstraction layers which allows very simple implementations.
 - ▶ **Keras** is very widely used and makes accessing TensorFlow very easy.
 - ▶ **PyTorch** is already conceptually a “high level” implementation.
- ▶ **Keras** can use various **backends** (implementations):
 - ▶ **TensorFlow**
 - ▶ **MXNet**
 - ▶ **Theano** is a pure python library for a wide class of array computation, not just Neural Networks. It was forked into **Aesara**...
 - ▶ **Microsoft Cognitive Toolkit**, but this is no longer in active development.
- ▶ See **Tensorflow or keras?**

Practical advice

- ▶ Explore recommendations. e.g. **Practical Advice for Building Deep Neural Networks**:
- ▶ As a starting point:
 - ▶ Use the “adam” optimizer
 - ▶ Use a ReLU activation function
 - ▶ Remember not to use an activation function for the output layer (except for classification, when use a sigmoid)
 - ▶ Add bias to every layer (shouldn’t have to worry about this in keras)
 - ▶ Whiten (normalize) your input data (we’ll see this in the workshop)
- ▶ **Don’t believe me.** Get other opinions, and try things yourself.

Debugging


- ▶ Check the **input data**...
- ▶ For many tasks:
 - ▶ **OVERFIT**. “Accuracy should be essentially 100% or 99.99%”.
If it isn't, the network isn't flexible enough, or learning correctly.
- ▶ Change the learning rate
- ▶ Decrease mini-batch size
- ▶ Remove batch normalization (this exposes NA values)
- ▶ Reconsider the architecture
- ▶ PLOT your results! training loss by epoch is a natural plot

Additional notes on learning

- ▶ Learning a Neural Network is still non-trivial. Start with this advice⁴
 - ▶ **Second order methods** are often used later in the fitting process, closer to the global optima.
 - ▶ **Hyperparameters** matter. Some optimisers, e.g. Adam, can tune them semi-automatically. Standard ones require **manual tuning** for e.g. step size.
- ▶ There is nothing here to prevent **overfitting**!

⁴Bengio 2012 **Practical Recommendations for Gradient-Based Training of Deep Architectures**

Learning rates



- ▶ **not** specific to neural networks
- ▶ But particularly important due to NN flexibility

Hints on overfitting

- ▶ Many optimizers include options for these tricks and more:
- ▶ **Penalize** large weights:
 - ▶ Ridge (L2) penalisation: $L = L_0 + \lambda \sum_{i,j} |W_{ij}|^2$
 - ▶ Lasso (L1) penalisation: $L = L_0 + \lambda \sum_{i,j} |W_{ij}|$
- ▶ **Dropout**:
 - ▶ New hyperparameter p_k for layer k : the **dropout rate**
 - ▶ Each learning step, with independently randomly set all outputs from a neuron to 0
- ▶ **Early stopping**:
 - ▶ retain a test dataset (from the training dataset)
 - ▶ evaluate performance on the held-out set
 - ▶ stop when this no longer increases

Further reading

- ▶ [Keras and PyTorch](#)
- ▶ [Tensorflow or keras?](#)
- ▶ [A performance focussed comparison: TensorFlow, PyTorch or MXNet?](#)
- ▶ [Tensorboard](#)
- ▶ [Brilliant.org on Backpropagation](#)
- ▶ [Practical Advice for Building Deep Neural Networks](#)