

Practicalities of Neural Nets

Daniel Lawson — University of Bristol

Lecture 07.2 (v2.0.0)

Implementing Neural Networks

- ▶ Implementations are best thought of in two classes.
- ▶ **Simple networks** have a restricted architecture and can be deployed “out of the box” as a Machine Learning tool.
 - ▶ Examples include `sklearn.linear_model.Perceptron`, R’s `neuralnet` packages, etc
 - ▶ Often either shallow or very simple hidden layer structure
- ▶ **Deep networks** require a complex specification of architecture and significant computational optimisation, so are very large (and mercifully, open source) endeavours
 - ▶ This is the focus here.

Deep NN Implementations

- ▶ There are two main libraries for deep neural networks:
- ▶ **TensorFlow**, developed by Google Brain.
 - ▶ Well documented
 - ▶ Easier to use
 - ▶ Industry standard
 - ▶ **Tensorboard** visualisation is useful
- ▶ **PyTorch**, developed by Facebook.
 - ▶ Newer, less support
 - ▶ Dynamical coding paradigm: graph can remodel in the light of the data
 - ▶ Debugging is easier? As the code is compiled at runtime, like native python

Using implementations

- ▶ **Tensorflow** is a low-level language. You can interact with it through abstraction layers which allows very simple implementations.
 - ▶ **Keras** is very widely used and makes accessing TensorFlow very easy.
 - ▶ **PyTorch** is already conceptually a “high level” implementation.
- ▶ **Keras** can use various **backends** (implementations):
 - ▶ **TensorFlow**
 - ▶ **MXNet**
 - ▶ **Theano** is a pure python library for a wide class of array computation, not just Neural Networks. It was forked into **Aesara**...
 - ▶ **Microsoft Cognitive Toolkit**, but this is no longer in active development.
- ▶ See **Tensorflow or keras?**

Practical advice

- ▶ Explore recommendations. e.g. **Practical Advice for Building Deep Neural Networks**:
- ▶ As a starting point:
 - ▶ Use the “adam” optimizer
 - ▶ Use a ReLU activation function
 - ▶ Remember not to use an activation function for the output layer (except for classification, when use a sigmoid)
 - ▶ Add bias to every layer (shouldn’t have to worry about this in keras)
 - ▶ Whiten (normalize) your input data (we’ll see this in the workshop)
- ▶ **Don’t believe me.** Get other opinions, and try things yourself.

Debugging


- ▶ Check the **input data**...
- ▶ For many tasks:
 - ▶ **OVERFIT**. “Accuracy should be essentially 100% or 99.99%”.
If it isn't, the network isn't flexible enough, or learning correctly.
- ▶ Change the learning rate
- ▶ Decrease mini-batch size
- ▶ Remove batch normalization (this exposes NA values)
- ▶ Reconsider the architecture
- ▶ PLOT your results! training loss by epoch is a natural plot

Additional notes on learning

- ▶ Learning a Neural Network is still non-trivial. Start with this advice¹
 - ▶ **Second order methods** are often used later in the fitting process, closer to the global optima.
 - ▶ **Hyperparameters** matter. Some optimisers, e.g. Adam, can tune them semi-automatically. Standard ones require **manual tuning** for e.g. step size.
- ▶ There is nothing here to prevent **overfitting**!

¹Bengio 2012 **Practical Recommendations for Gradient-Based Training of Deep Architectures**

Learning rates



- ▶ **not** specific to neural networks
- ▶ But particularly important due to NN flexibility

Hints on overfitting

- ▶ Many optimizers include options for these tricks and more:
- ▶ **Penalize** large weights:
 - ▶ Ridge (L2) penalisation: $L = L_0 + \lambda \sum_{i,j} |W_{ij}|^2$
 - ▶ Lasso (L1) penalisation: $L = L_0 + \lambda \sum_{i,j} |W_{ij}|$
- ▶ **Dropout**:
 - ▶ New hyperparameter p_k for layer k : the **dropout rate**
 - ▶ Each learning step, with independently randomly set all outputs from a neuron to 0
- ▶ **Early stopping**:
 - ▶ retain a test dataset (from the training dataset)
 - ▶ evaluate performance on the held-out set
 - ▶ stop when this no longer increases

Reflection

- ▶ Understand the tools available for neural networks
- ▶ Be able to use high-level implementations efficiently

Further reading

- ▶ [Keras and PyTorch](#)
- ▶ [Tensorflow or keras?](#)
- ▶ [A performance focussed comparison: TensorFlow, PyTorch or MXNet?](#)
- ▶ [Tensorboard](#)
- ▶ [Brilliant.org on Backpropagation](#)
- ▶ [Practical Advice for Building Deep Neural Networks](#)